

Simple Linear Regression

Learning Activities

- learn to use scikit-learn to implement simple linear regression
- learn to create, train, and test a linear regression model on real data

Import needed packages

For this lab, you will need to have the following packages:

- NumPy
- Matplotlib
- Pandas
- Scikit-learn

Execute the below to ensure you have the above to avoid importing issues.

- !pip install numpy==2.2.0
- !pip install pandas==2.2.3
- !pip install scikit-learn==1.6.0
- !pip install matplotlib==3.9.3

```
In [1]: !pip install -q numpy pandas scikit-learn matplotlib
```

```
In [2]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
%matplotlib inline
```

Load the data

Use the pandas library to load the dataset

```
In [3]: df = pd.read_csv('my-fuel-consumption.csv')
```

```
In [4]: df.sample(5)
```

Out[4]:

	Model year	Make	Model	Vehicle class	Engine size (L)	Cylinders	Transmission	Fuel type	City (L/100 km)	Highway (L/100 km)	Combined (L/100 km)	Combined (mpg)	CO2 emissions (g/km)	CO2 rating	Smog rating
380	2026	Toyota	Corolla (3-mode)	Compact	2.0	4	AV10	X	7.6	5.9	6.8	42	160	6	5
390	2026	Toyota	Crown Signia AWD	Station wagon: Small	2.5	4	AV	X	6.1	6.3	6.2	46	144	7	6
202	2026	Hyundai	Elantra	Mid-size	2.0	4	AV1	X	7.8	5.9	6.9	41	162	6	6
192	2026	GMC	Yukon	Sport utility vehicle: Standard	5.3	8	A10	X	15.7	12.0	14.0	20	329	3	6
4	2026	Audi	Q8 55 TFSI quattro	Sport utility vehicle: Standard	3.0	6	AS8	Z	13.6	10.4	12.1	23	285	4	4

Understand the data

FuelConsumption.csv:

We will use a fuel consumption dataset, `FuelConsumption.csv`, which contains model-specific ratings and estimated CO2 emissions for new vehicles for retail sale in Canada.

Features:

- MODEL YEAR, MAKE, MODEL

- VEHICLE CLASS, ENGINE SIZE
- CYLINDERS, TRANSMISSION, FUEL TYPE
- FUEL CONSUMPTION CITY (L/100 km), FUEL CONSUMPTION HWY (L/100 km)
- FUEL CONSUMPTION COMBINED (L/100 km), FUEL CONSUMPTION COMBINED MPG (MPG)
- CO2 EMISSIONS (g/km)

Activity: Create a simple linear regression model from one of these features to predict CO2 emissions of unobserved cars based on that feature.

Basic data exploration

In [5]: `df.describe()`

Out[5]:	Model year	Engine size (L)	Cylinders	City (L/100 km)	Highway (L/100 km)	Combined (L/100 km)	Combined (mpg)	CO2 emissions (g/km)	CO2 rating	Smog rating
count	437.0	437.000000	437.000000	437.000000	437.000000	437.000000	437.000000	437.000000	437.000000	437.000000
mean	2026.0	3.051945	5.427918	12.097712	9.372311	10.873227	28.082380	255.219680	4.409611	4.956522
std	0.0	1.305678	1.876325	3.495363	2.308001	2.896746	8.522371	67.033008	1.340683	1.371544
min	2026.0	1.200000	3.000000	4.400000	4.400000	4.400000	15.000000	104.000000	1.000000	2.000000
25%	2026.0	2.000000	4.000000	9.600000	7.600000	8.800000	22.000000	204.000000	4.000000	4.000000
50%	2026.0	2.700000	4.000000	12.000000	9.300000	10.700000	26.000000	259.000000	4.000000	5.000000
75%	2026.0	3.500000	6.000000	14.600000	10.700000	12.700000	32.000000	298.000000	5.000000	6.000000
max	2026.0	6.700000	12.000000	23.500000	16.600000	19.200000	64.000000	451.000000	8.000000	7.000000

Engine size (L)

- **Stats:** mean = 3.05, std = 1.31, min = 1.2, 25% = 2.0, median = 2.7, 75% = 3.5, max = 6.7
- **Meaning:**
 - Average engine size is about 3 liters.
 - Most engines range between 2 and 3.5 liters.
 - Most cars are “medium-sized” in engine power.

Cylinder

- **Stats:** mean = 5.43, min = 3, 25% = 4, median = 4, 75% = 6, max = 12
- **Meaning:**
 - Most cars have 4–6 cylinders.
 - Some small cars have 3 cylinders, and a few large engines have 12.
 - Think of cylinders as “engine muscles”—most cars have medium strength.

City fuel consumption (L/100 km)

- **Stats:** mean = 12.1, min = 4.4, 25% = 9.6, median = 12, 75% = 14.6, max = 23.5
- **Meaning:**
 - On average, cars use 12 L of fuel per 100 km in the city.
 - Some are very efficient (4.4 L), while others are thirsty (23.5 L).
 - Most cars use fuel moderately; city driving can be expensive for less efficient vehicles.

Combined fuel consumption (L/100 km)

- **Stats:** mean = 10.87, min = 4.4, 25% = 8.8, median = 10.7, 75% = 12.7, max = 19.2
- **Meaning:**
 - Average fuel consumption combining city and highway driving is ~10.9 L/100 km.
 - This gives a practical view of “typical” driving fuel use.

Combined fuel efficiency (mpg)

- **Stats:** mean = 28.08, min = 15, 25% = 22, median = 26, 75% = 32, max = 64
- **Meaning:**
 - Cars get about 28 miles per gallon on average.
 - Lower mpg = more frequent trips to the gas station; higher mpg = fewer stops.

CO₂ emissions (g/km)

- **Stats:** mean = 255, min = 104, 25% = 204, median = 259, 75% = 298, max = 451
- **Meaning:**
 - Average car emits 255 g/km of CO₂.

- Small cars can be very clean (104 g/km), while big engines pollute more (451 g/km).
- Most cars are moderately polluting; aiming below 204 g/km is greener.

In [6]: `df.isnull().sum()`

Out[6]:

Model year	0
Make	0
Model	0
Vehicle class	0
Engine size (L)	0
Cylinders	0
Transmission	0
Fuel type	0
City (L/100 km)	0
Highway (L/100 km)	0
Combined (L/100 km)	0
Combined (mpg)	0
CO2 emissions (g/km)	0
CO2 rating	0
Smog rating	0
dtype: int64	

Missing values check

- All columns have 437 valid entries
- This means there are no missing numbers in any of these features

Feature Selection

Select features that might be indicative of CO2 emission

In [7]: `df.rename(columns={
 "Model year": "ModelYear", "Make": "Make", "Model": "Model", "Vehicle class": "VehicleClass",
 "Engine size (L)": "EngineSize", "Cylinders": "Cylinders", "Transmission": "Transmission",
 "Fuel type": "FuelType", "City (L/100 km)": "Fuel_CityL100Km", "Fuel_Highway (L/100 km)": "HighwayL100Km",
 "Combined (L/100 km)": "FuelConsumption_Comb", "Combined (mpg)": "FuelConsumption_Mpg",
 "CO2 emissions (g/km)": "CO2Emissions", "CO2 rating": "CO2Rating", "Smog rating": "SmogRating"
}, inplace=True)`

```
In [8]: cdf = df[['EngineSize', 'Cylinders', 'FuelConsumption_Comb', 'CO2Emissions']]  
cdf.sample(4)
```

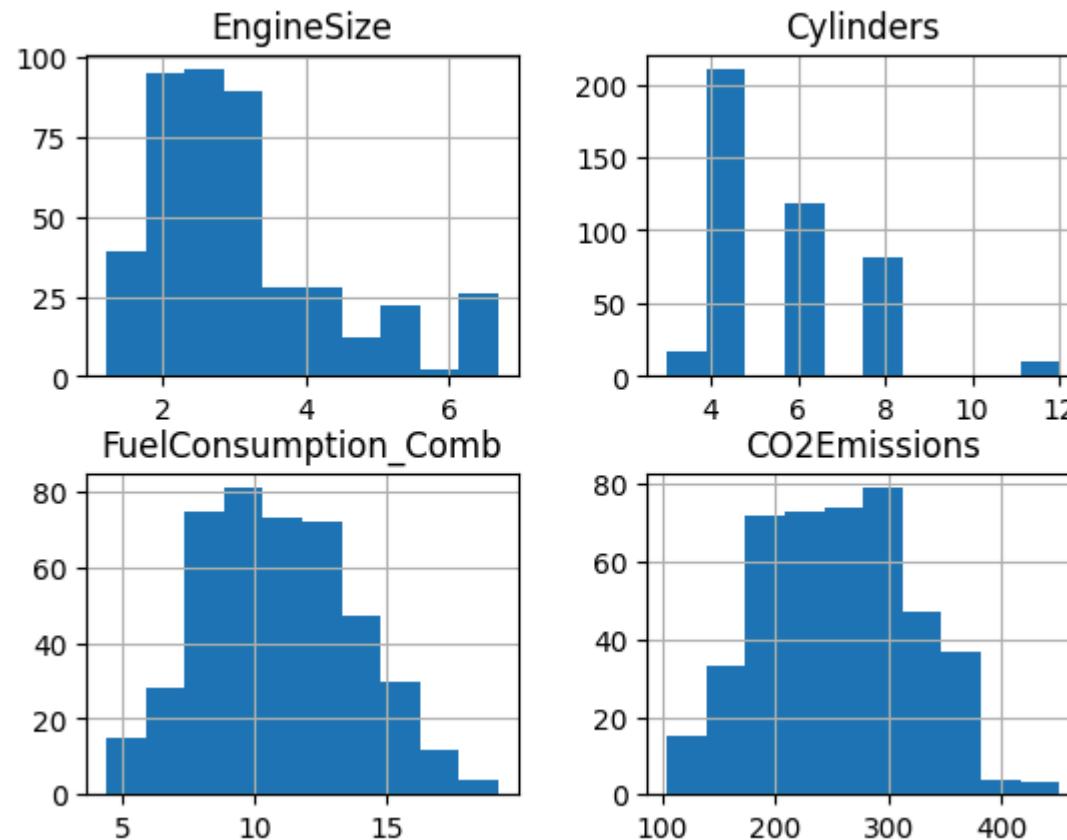
```
Out[8]:
```

	EngineSize	Cylinders	FuelConsumption_Comb	CO2Emissions
214	2.5	4	8.1	189
317	2.0	4	7.4	171
391	1.6	3	10.5	245
285	2.5	4	6.0	141

Feature Visualization

Histogram, scatter plot

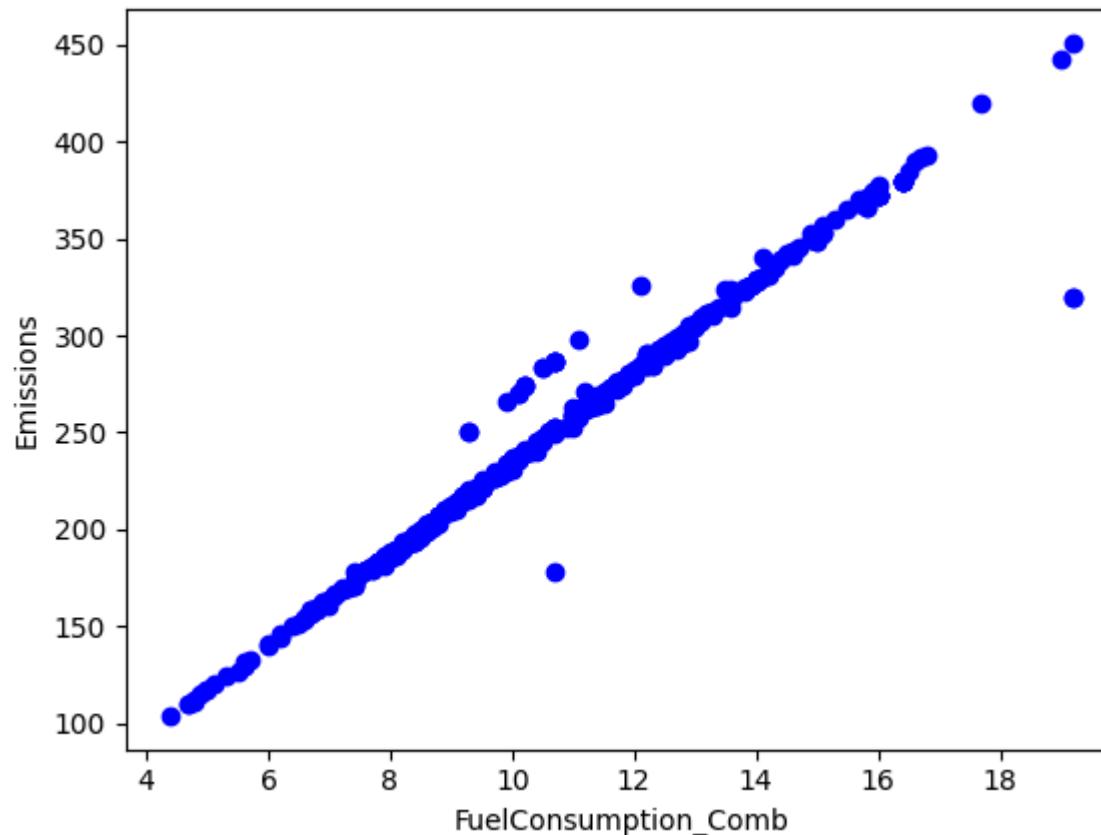
```
In [9]: viz = cdf[['EngineSize','Cylinders', 'FuelConsumption_Comb', 'CO2Emissions']]  
viz.hist()  
plt.show()
```



Insights

- Most cars have 4,6,8 cylinders
- Engine sizes are between 2 and 4

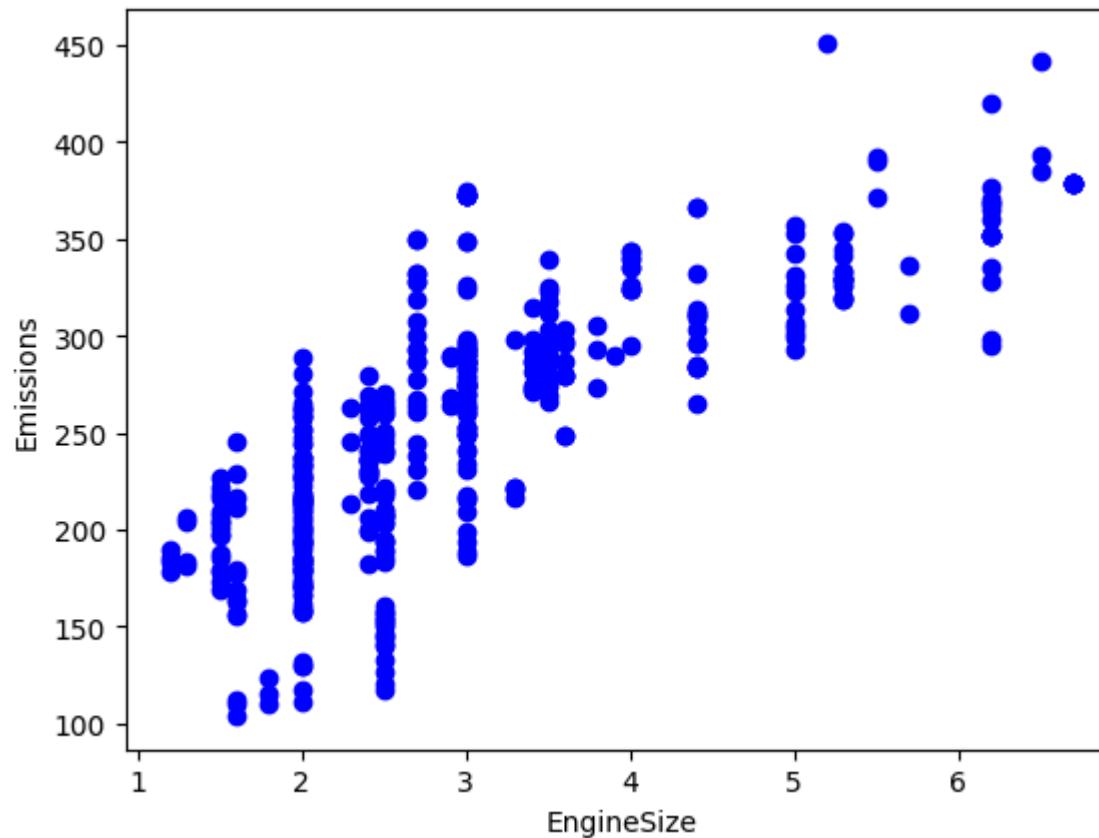
```
In [10]: plt.scatter(cdf.FuelConsumption_Comb, cdf.CO2Emissions, color='blue')
plt.xlabel("FuelConsumption_Comb")
plt.ylabel("Emissions")
plt.show()
```



Insights

- One car group has a strong linear relationship between it's combined fuel consumption and CO2 emissions.
- Another group has a linear relation but it is weak.

```
In [11]: plt.scatter(cdf.EngineSize, cdf.CO2Emissions, color='blue')
plt.xlabel("EngineSize")
plt.ylabel("Emissions")
plt.show()
```



Insights

Positive Correlation

- As engine size increases, CO₂ emissions also increase.
- Upward trend from left (small engines) to right (large engines).

Smaller Engines - Lower Emissions

- Vehicles with 1.5L–2.5L engines mostly emit between 100–250 g/km.
- These are usually compact cars or efficient models.

Data Spread Shows Some Variation

- Even within the same engine size:
- Emissions vary slightly depending on vehicle design, fuel type

Extract the input features and labels from the dataset

- Let us pick engine size
- Extract the input feature and target / output variables.

In [15]:

```
#For EngineSize
X = cdf.EngineSize.to_numpy()
y = cdf.CO2Emissions.to_numpy()
```

Create, Train and Test Datasets

Steps

- Split the dataset into two parts: training set and testing set.
- Train a SLM using the training set.
- Use the trained model to make predictions on the unseen testing set.
- The test data already has the real answers.
- After training the model, use the test data to check its predictions.
- Compare the predicted values with the actual values.
- This tells how well the model performs on new, unseen data.
- This measurement is called **out-of-sample accuracy**.

In [16]:

```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.2, random_state=42)
```

In [17]:

```
type(X_train), np.shape(X_train), np.shape(X_train)
```

Out[17]:

```
(numpy.ndarray, (349,), (349,))
```

Build a Simple Linear Regression Model

```
In [19]: #Using scikit-Learn
from sklearn import linear_model

#create a model
regress = linear_model.LinearRegression()

#train the model on training data
regress.fit(X_train.reshape(-1,1), y_train)

#print the coefficients
print('Coefficients:', regress.coef_[0])
print('Intercept:', regress.intercept_)
```

Coefficients: 39.886630889809524
Intercept: 134.94676409410977

Coefficient and Intercept: They define the slope and best fit line to the training data

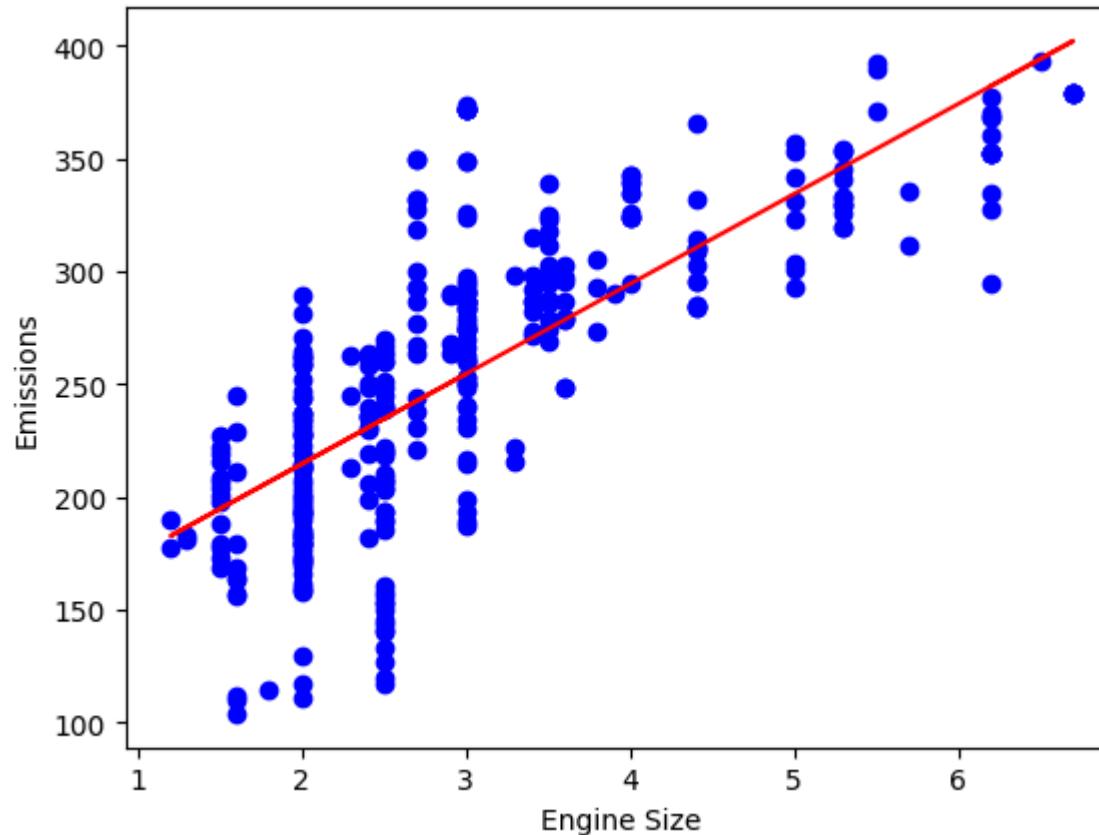
Visualize Model Outputs

Plot Goodness-of-fit of the model to the training data.

The regression model is the line given by $y = \text{intercept} + \text{coefficient} \cdot x$.

```
In [20]: plt.scatter(X_train, y_train, color='blue')
plt.plot(X_train, regress.coef_*X_train + regress.intercept_, '-r')
plt.xlabel("Engine Size")
plt.ylabel("Emissions")
```

```
Out[20]: Text(0, 0.5, 'Emissions')
```



Model Evaluation

Compare the actual values and predicted values to calculate the accuracy of the regression model.

- **Mean Absolute Error (MAE)**: Average of the absolute differences between actual and predicted values.
- **Mean Squared Error (MSE)**: Average of squared differences between actual and predicted values; measures model error.
- **Root Mean Squared Error (RMSE)**: Square root of MSE; in the same units as the data, easier to interpret.
- **R-squared (R^2)**: Measures how well the model fits the data; closer to 1 means better fit.

```
In [24]: from sklearn.metrics import mean_absolute_error, mean_squared_error, root_mean_squared_error, r2_score  
  
#Use the predict method to make test prediction  
y_test_ = regress.predict(X_test.reshape(-1,1))
```

```
print("Mean absolute error: %.2f" % mean_absolute_error(y_test_, y_test))
print("Mean squared error: %.2f" % mean_squared_error(y_test_, y_test))
print("Root mean squared error: %.2f" % root_mean_squared_error(y_test_, y_test))
print("R2-score: %.2f" % r2_score(y_test_, y_test))
```

Mean absolute error: 30.08
Mean squared error: 1583.53
Root mean squared error: 39.79
R2-score: 0.53

- **Mean Absolute Error (MAE = 30.08):** On average, the model's predicted CO₂ emissions are about 30 g/km off from the actual values.
- **Mean Squared Error (MSE = 1583.53):** Average of squared errors (g/km)²; larger mistakes contribute more, highlighting big prediction errors.
- **Root Mean Squared Error (RMSE = 39.79):** Square root of MSE; gives average deviation in original units, so predictions are about 39.79 g/km off on average.
- **R²-score (0.53):** The model explains 53% of the variance in CO₂ emissions, indicating a moderate fit.

Summary: The model moderately fits the CO₂ emissions data, capturing some patterns but with typical prediction errors around 40 g/km. Large deviations have a bigger impact due to squaring in MSE.